

Week-5

Result:-

The screenshot shows two Jupyter Notebook sessions. The top session displays the model summary, and the bottom session shows the code used to build and train the model.

Model Summary:

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 28, 28, 32)	2,432
max_pooling2d_8 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_9 (Conv2D)	(None, 10, 10, 64)	51,264
max_pooling2d_9 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_4 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 64)	102,464
dense_11 (Dense)	(None, 10)	650

Total params: 156,810 (612.54 KB)
Trainable params: 156,810 (612.54 KB)
Non-trainable params: 0 (0.00 B)

Code Implementation:

```
[47]:  
import tensorflow as tf  
from tensorflow.keras import datasets, layers, models  
from tensorflow.keras.losses import MeanAbsoluteError  
(x_train, y_train), (x_test, y_test) = datasets.cifar10.load_data()  
x_train, x_test = x_train / 255.0, x_test / 255.0  
model = models.Sequential([  
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),  
    layers.MaxPooling2D((2, 2)),  
    layers.Conv2D(64, (3, 3), activation='relu'),  
    layers.MaxPooling2D((2, 2)),  
    layers.Flatten(),  
    layers.Dense(64, activation='relu'),  
    layers.Dense(10)  
)  
model.compile(optimizer='adam',  
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              metrics=['mae', 'accuracy'])  
model.summary()  
history = model.fit(x_train, y_train, epochs=5, batch_size=50,  
                     validation_data=(x_test, y_test))  
test_loss, test_mae, test_acc = model.evaluate(x_test, y_test, verbose=2)  
print(f"\nTest MAE: {test_mae:.4f}")  
print(f"Test Accuracy: {test_acc:.4f}")
```

The screenshot shows a Jupyter Notebook running on localhost:8889. The notebook title is "Week-5-CNN_EUR-USD_ML_in-Finance_Final.ipynb". The code cell displays training metrics for a CNN model. The output shows:

```
Trainable params: 156,810 (612.54 KB)
Non-trainable params: 0 (0.00 B)

Epoch 1/5
1000/1000    45s 40ms/step - accuracy: 0.4532 - loss: 1.5154 - mae: 4.8374 - val_accuracy: 0.5308 - val_loss: 1.2911 - val_mae: 4.8413
Epoch 2/5
1000/1000    41s 41ms/step - accuracy: 0.5859 - loss: 1.1718 - mae: 4.9665 - val_accuracy: 0.5840 - val_loss: 1.1540 - val_mae: 5.0391
Epoch 3/5
1000/1000    39s 39ms/step - accuracy: 0.6334 - loss: 1.0423 - mae: 5.1280 - val_accuracy: 0.6074 - val_loss: 1.1110 - val_mae: 5.2448
Epoch 4/5
1000/1000    39s 39ms/step - accuracy: 0.6657 - loss: 0.9512 - mae: 5.2313 - val_accuracy: 0.6407 - val_loss: 1.0253 - val_mae: 5.4532
Epoch 5/5
1000/1000    39s 38ms/step - accuracy: 0.6888 - loss: 0.8872 - mae: 5.3908 - val_accuracy: 0.6616 - val_loss: 0.9743 - val_mae: 5.1922
313/313 - 5s - 16ms/step - accuracy: 0.6616 - loss: 0.9743 - mae: 5.1922

Test MAE: 5.1922
Test Accuracy: 0.6616
```

Below the code cell, there are two empty input fields labeled "[]:" followed by a placeholder "Click to add a cell." The system tray at the bottom right shows the date and time as 26-10-2025, 23:07:47.